

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/382494872>

Llama 3.1: An In-Depth Analysis of the Next Generation Large Language Model

Preprint · July 2024

DOI: 10.13140/RG.2.2.10628.74882

CITATIONS

0

READS

6,366

2 authors:



[Kira Sam](#)

Università del Salento

16 PUBLICATIONS 3 CITATIONS

[SEE PROFILE](#)



[Raja Vavekanand](#)

Benazir Bhutto Shaheed University Lyari (BBSUL) Karachi

30 PUBLICATIONS 63 CITATIONS

[SEE PROFILE](#)

Llama 3.1: An In-Depth Analysis of the Next-Generation Large Language Model

Raja Vavekanand*, Kira Sam*

*Datalink Research and Technology Lab

Abstract

Llama 3.1 405B, the first open model with 405 billion parameters, rivals top AI models in knowledge, math, tool use, and multilingual translation. Its release introduces opportunities for innovation, including synthetic data generation and model distillation. Upgraded 8B and 70B models feature a 128K context length and enhanced reasoning. Evaluated on over 150 benchmarks, Llama 3.1 excels in diverse tasks and competes with leading models like GPT-4. Openly available, it enables developers to customize and deploy advanced AI solutions, fostering broad-based AI innovation and application.

Keywords: *Llama 3.1 405B, AI model, multilingual translation, synthetic data generation, model distillation, open source.*

1. Introduction

In the rapidly evolving field of artificial intelligence (AI), foundation models have emerged as a crucial component, serving as versatile tools for various modalities such as language, vision, speech, and more. These models underpin many contemporary AI systems, supporting many tasks. The development of these models involves two primary stages: pre-training and post-training. Pre-training involves massive-scale training using tasks like next-word prediction or captioning while post-training refines the model to follow instructions, align with human preferences, and enhance specific capabilities like coding and reasoning. Llama 3, is a new suite of foundation models for language, which we refer to as the Llama 3 Herd. This suite supports multilingual, coding, reasoning, and tool usage, making it a powerful tool for various AI applications. The flagship model is a dense Transformer with 405 billion parameters, capable of processing information in a context window of up to 128K tokens. The Llama 3 models include 8B, 70B, and 405B parameter versions, each tailored for various AI tasks.

1.1. Key Factors

Three key factors drive the development of high-quality foundation models: data, scale, and managing complexity. We optimized these factors in Llama 3's development:

- **Data:** We significantly enhanced the quantity and quality of data used for pre-training and post-training. This involved improved pre-processing and curation pipelines and rigorous quality assurance and filtering approaches. Llama 3 was pre-trained on approximately 15 trillion multilingual tokens, compared to 1.8 trillion tokens for Llama 2.
- **Scale:** We trained our models at an unprecedented scale. The flagship model, with 405 billion trainable parameters, was pre-trained using 3.8×10^{25} FLOPs, nearly 50 times more than the largest version of Llama 2. This model was trained on 15.6 trillion text tokens. According to scaling laws, our flagship model outperforms smaller models trained using the same procedures.
- **Managing Complexity:** We made design choices to maximize scalability. For instance, we opted for a standard dense Transformer model architecture with minor adaptations,

prioritizing training stability over complex mixture-of-experts models. Our post-training procedure is relatively simple, relying on supervised fine-tuning (SFT), rejection sampling (RS), and direct preference optimization (DPO), avoiding more complex and less stable reinforcement learning algorithms.

Llama 3, a new AI model, is set to revolutionize the field of AI research and application. It comprises three multilingual language models with 8B, 70B, and 405B parameters, each with its unique performance. The flagship model, Llama 3.1 405B, outperforms leading models like GPT-4 in various tasks and offers a better balance between helpfulness and harmlessness. The models will be released under an updated Llama 3 Community License, aiming to foster innovation and accelerate progress towards artificial general intelligence. The open release includes pre-trained and post-trained versions of the 405B parameter model and a new version of the Llama Guard model for input and output safety. The development of multimodal extensions to the Llama 3 models, which support image recognition, video recognition, and speech understanding, is still in development. These extensions will further enhance the capabilities of Llama 3, making it an even more powerful tool for various AI applications. Llama 3 represents a significant advancement in the field of AI research and application. Its improved performance, efficiency, and scalability make it an ideal tool for various AI tasks. We believe that Llama 3 will revolutionize the field of AI and accelerate progress towards artificial general intelligence.

Category Benchmark	Llama 3.1 405B	Nemotron 4 340B Instruct	GPT-4 (0125)	GPT-4 Omni	Claude 3.5 Sonnet
General					
MMLU (0-shot, CoT)	88.6	78.7 (non-CoT)	85.4	88.7	88.3
MMLU PRO (5-shot, CoT)	73.3	62.7	64.8	74.0	77.0
IFEval	88.6	85.1	84.3	85.6	88.0
Code					
HumanEval (0-shot)	89.0	73.2	86.6	90.2	92.0
MBPP EvalPlus (base) (0-shot)	88.6	72.8	83.6	87.8	90.5
Math					
GSM8K (8-shot, CoT)	96.8	92.3 (0-shot)	94.2	96.1	96.4 (0-shot)
MATH (0-shot, CoT)	73.8	41.1	64.5	76.6	71.1
Reasoning					
ARC Challenge (0-shot)	96.9	94.6	96.4	96.7	96.7
GPQA (0-shot, CoT)	51.1	-	41.4	53.6	59.4
Tool use					
BFCL	88.5	86.5	88.3	80.5	90.2
Nexus	58.7	-	50.3	56.1	45.7
Long context					
ZeroSCROLLS/QuALITY	95.2	-	95.2	90.5	90.5
InfiniteBench/En.MC	83.4	-	72.1	82.5	-
NIH/Multi-needle	98.1	-	100.0	100.0	90.8
Multilingual					
Multilingual MGSM (0-shot)	91.6	-	85.9	90.5	91.6

Category Benchmark	Llama 3.1 8B	Gemma 2 9B IT	Mistral 7B Instruct	Llama 3.1 70B	Mistral 8x22B Instruct	GPT 3.5 Turbo
General						
MMLU (0-shot, CoT)	73.0	72.3 (5-shot, non-CoT)	60.5	86.0	79.9	69.8
MMLU PRO (5-shot, CoT)	48.3	-	36.9	66.4	56.3	49.2
IFEval	80.4	73.6	57.6	87.5	72.7	69.9
Code						
HumanEval (0-shot)	72.6	54.3	40.2	80.5	75.6	68.0
MBPP EvalPlus (base) (0-shot)	72.8	71.7	49.5	86.0	78.6	82.0
Math						
GSM8K (8-shot, CoT)	84.5	76.7	53.2	95.1	88.2	81.6
MATH (0-shot, CoT)	51.9	44.3	13.0	68.0	54.1	43.1
Reasoning						
ARC Challenge (0-shot)	83.4	87.6	74.2	94.8	88.7	83.7
GPQA (0-shot, CoT)	32.8	-	28.8	46.7	33.3	30.8
Tool use						
BFCL	76.1	-	60.4	84.8	-	85.9
Nexus	38.5	30.0	24.7	56.7	48.5	37.2
Long context						
ZeroSCROLLS/QuALITY	81.0	-	-	90.5	-	-
InfiniteBench/En.MC	65.1	-	-	78.2	-	-
NIH/Multi-needle	98.8	-	-	97.5	-	-
Multilingual						
Multilingual MGSM (0-shot)	68.9	53.2	29.9	86.9	71.1	51.4

Figure 1. The Llama 3.1 model was evaluated on 150 benchmark datasets across various languages and human evaluations. Results showed it competes with leading foundation models like GPT-4, GPT-4o, and Claude 3.5 Sonnet, and smaller models with closed and open models. (Meta,2024)

2. Overview

In the rapidly evolving field of artificial intelligence (AI), large language models (LLMs) have emerged as a crucial component, transforming the way we interact with technology. Llama 3, a new AI model, is set to revolutionize the field of AI research and application. It comprises three multilingual language models with 8B, 70B, and 405B parameters, each with its unique performance. The flagship model, Llama 3.1 405B, outperforms leading models like GPT-4 in various tasks and offers a better balance between helpfulness and harmlessness. The models will be released under an updated Llama 3 Community License, aiming to foster innovation and accelerate progress towards artificial general intelligence. The open release also includes pre-trained and post-trained versions of the 405B parameter model and a new version of the Llama Guard model for input and output safety. The development of multimodal extensions to the Llama 3 models, which support image recognition, video recognition, and speech understanding, is still in development. Meta AI's Llama 3.1 is a state-of-the-art LLM that has set new standards in the industry. As part of the Llama 3 suite of foundation models, Llama 3.1 is designed to support multilingualism, coding, reasoning, and tool usage, making it an incredibly versatile tool.

Llama 3.1 boasts an impressive array of features that make it one of the most powerful LLMs available:

- **405 billion parameters:** Llama 3.1 has an unprecedented 405 billion trainable parameters, allowing it to process and generate text with unparalleled accuracy and context.
- **Multilingual Support:** Llama 3.1 can process and generate text in multiple languages, including but not limited to English, Spanish, French, German, Italian, Portuguese, Dutch, Russian, Chinese, Japanese, and Korean, making it a valuable tool for global communication.
- **Context window of up to 128K tokens:** Llama 3.1 can process long-range dependencies and context, making it suitable for tasks like document summarization, question answering, and text generation.
- **Improved performance and efficiency:** Llama 3.1 outperforms its predecessor, Llama 2, and other leading models like GPT-4 in various tasks while offering better efficiency and scalability.
- **Simple and stable architecture:** Llama 3.1 uses a standard dense Transformer model architecture with minor adaptations, prioritizing training stability and scalability.




Figure 2. Llama 3.1 Model Evaluation. (Meta,2024)

3. Architecture

Llama 3.1 405B is one of the largest LLMs available, following the transformer architecture. It consists of an encoder and a decoder, processing input text and generating output text. To train on over 15 trillion tokens, the model was optimized and trained on over 16 thousand H100 GPUs. To keep the development process scalable and straightforward, the model was designed with a standard decoder-only transformer model architecture and an iterative post-training procedure. This allowed for the creation of high-quality synthetic data for each round and improved performance. Compared to previous versions of Llama, improvements were made to the quantity and quality of data used for pre-and post-training. These included more careful pre-processing and curation pipelines, rigorous quality assurance, and filtering approaches for post-training data. The new flagship model outperforms smaller models trained using the same procedure, and the 405B parameter model was used to improve post-training quality. To support large-scale production inference, the models were quantized from 16-bit to 8-bit numerics, lowering compute requirements and allowing the model to run within a single server node.




Figure 3. Illustration of the overall architecture and training. Llama 3.1 405B, the largest model yet, was trained on 15 trillion tokens using over 16,000 H100 GPUs, overcoming significant challenges in training. (Meta,2024)

4. Training Data

Llama 3.1, the latest generation of foundation models, was trained on an unprecedented scale, utilizing a dataset comprising 15 trillion tokens. This extensive training dataset, sourced from a variety of high-quality and diverse text sources, plays a crucial role in the model's ability to generate coherent, contextually relevant, and human-like text across different languages and tasks. The diverse nature of the training data ensures that Llama 3.1 can understand and respond to a wide range of inputs, making it a versatile tool for numerous applications.

4.1. Sources of Training Data

The training data for Llama 3.1 was gathered from multiple sources to cover a broad spectrum of knowledge and linguistic styles. Key sources include:

4.1.1. Web Pages

- **Variety and Breadth:** Web pages offer a vast array of content, from blog posts and forums to news articles and educational resources. This variety ensures the model is exposed to different writing styles, terminologies, and subject matters.
- **Real-World Language Usage:** The inclusion of web pages allows the model to learn from real-world language use, capturing colloquial expressions, idioms, and the informal language often used online.

4.1.2. Books

- **Depth of Knowledge:** Books provide in-depth coverage of topics, offering structured and detailed content. This helps the model understand complex ideas and concepts.
- **Literary Quality:** The high literary quality of books aids in training the model to generate well-structured and grammatically correct sentences.

4.1.3. Articles

- **Topical and Current Information:** Articles, including those from journals, magazines, and online publications, offer up-to-date information on various subjects, from scientific discoveries to current events.
- **Balanced Perspectives:** Articles often present multiple viewpoints on an issue, helping the model develop a nuanced understanding of controversial topics.

4.1.4. Research Papers

- **Technical and Specialized Knowledge:** Research papers contribute specialized knowledge from various academic fields, enabling the model to generate and understand technical content.
- **Scientific Rigor:** The structured and precise language of research papers helps in training the model to produce accurate and reliable information.

4.2. Pre-Processing and Curation

The training data for Llama 3.1 underwent rigorous pre-processing and curation to ensure its quality and relevance. Key steps in this process include:

- **Data Cleaning:** Removing noise and irrelevant content from the raw data to ensure only useful information is retained.
- **Duplication Removal:** Identifying and removing duplicate entries to prevent overfitting and redundancy in the training data.
- **Quality Filtering:** Implementing filters to exclude low-quality or unreliable sources, ensuring the training data is of the highest possible standard.

4.3. Multilingual Training

Llama 3.1's training data includes content in multiple languages, enabling the model to understand and generate text in various linguistic contexts. This multilingual capability is essential for applications requiring language translation, cross-lingual information retrieval, and multilingual conversational agents.

4.4. Data Quantity and Scale

The sheer volume of 15 trillion tokens used in training Llama 3.1 is a testament to the model's scale and the computational resources involved. This extensive dataset allows the model to learn from a vast amount of information, enhancing its generalization capabilities and performance across different tasks.

4.5. Impact of Diverse Training Data

- **Enhanced Understanding:** Exposure to different content types and linguistic styles improves the model's ability to understand and generate diverse text.
- **Contextual Relevance:** The variety of sources helps the model provide contextually relevant responses, making it more useful in real-world applications.
- **Robustness:** A wide-ranging training dataset contributes to the model's robustness, enabling it to handle varied and unexpected inputs effectively.

4.6. Challenges and Considerations

Training a model like Llama 3.1 on such a vast and diverse dataset involves several challenges:

- **Data Quality Management:** Ensuring the quality of the training data is paramount. Poor-quality data can lead to inaccurate or biased outputs.
- **Computational Resources:** Training on 15 trillion tokens requires significant computational power and storage capacity, highlighting the importance of efficient resource management.
- **Ethical Considerations:** Care must be taken to address ethical issues related to data privacy, bias, and the responsible use of AI.

5. Capabilities

Llama 3.1, the latest iteration of the Llama model family, demonstrates remarkable proficiency in a range of natural language processing (NLP) tasks. Leveraging its extensive training on a massive and diverse dataset, Llama 3.1 excels in text generation, language translation, summarization, and question-answering. These capabilities make it a versatile tool for various applications, including chatbots, content creation, multilingual communication, information retrieval, and more.

5.1. Text Generation

Llama 3.1 excels in generating coherent, context-specific text, which is a fundamental requirement for numerous applications. Its ability to produce human-like responses makes it suitable for:

- **Chatbots:** The model can engage in natural and contextually relevant conversations, enhancing user interactions with virtual assistants and customer support systems.
- **Content Creation:** Llama 3.1 can assist writers by generating creative content, drafting articles, or providing suggestions for story development. This capability can significantly boost productivity in content marketing, blogging, and creative writing.

5.2. Language Translation

Llama 3.1's multilingual training enables it to perform high-accuracy translations between different languages. Key applications include:

- **Multilingual Communication:** The model can facilitate communication across language barriers, making it an invaluable tool for businesses operating in global markets.
- **Localization:** It can assist in localizing content, ensuring that products, services, and information are accessible to diverse audiences in their native languages.

5.3. Summarization

Summarizing lengthy texts into concise and meaningful summaries is another area where Llama 3.1 excels. This capability is particularly useful for:

- **Information Retrieval:** The model can provide quick summaries of long documents, making it easier for users to grasp essential information without reading through entire texts.
- **Content Digest:** It can generate summaries for news articles, research papers, and reports, helping professionals stay informed and make decisions based on condensed information.

5.4. Question Answering

Llama 3.1 can accurately answer questions based on input text, demonstrating a deep understanding of the context and content. This functionality is applicable to:

- **Customer Support:** The model can handle a wide range of customer inquiries by providing accurate and context-aware responses, thereby enhancing the efficiency of support systems.
- **Educational Tools:** It can assist students and educators by answering questions related to study materials, providing explanations, and offering additional insights.
- **Coding and reasoning:** Llama 3.1 can generate code, explain programming concepts, and perform logical reasoning, making it suitable for applications like code completion, debugging, and programming education.

Llama 3.1's capabilities in text generation, language translation, summarization, and question-answering underscore its versatility and utility across various domains. Its ability to understand and generate human-like text makes it a powerful tool for enhancing communication, improving productivity, and facilitating access to information. As AI technology continues to evolve, models like Llama 3.1 will play a crucial role in shaping the future of human-computer interaction and information processing.

6. Application

Llama 3.1, with its advanced capabilities in text generation, language translation, summarization, and question answering, opens up a wide array of applications across various

industries. Here, we delve into three primary applications: customer service, content creation, and language learning.

6.2. Customer Service

The integration of Llama 3.1 into customer service operations can revolutionize the way businesses interact with their customers. Key benefits include:

6.2. 24/7 Support

- **Always Available:** Chatbots powered by Llama 3.1 can provide round-the-clock support, ensuring that customer inquiries are addressed at any time of day. This constant availability helps in reducing response times and improving customer satisfaction.
- **Consistency:** The model delivers consistent and accurate responses, maintaining a uniform standard of service. This consistency is crucial for building trust and reliability in customer interactions.

6.3. Handling High Volume

- **Scalability:** Llama 3.1 can handle multiple queries simultaneously, making it possible to manage high volumes of customer inquiries without the need for a proportional increase in human support staff. This scalability is particularly beneficial during peak times or promotional events.
- **Efficiency:** By automating routine inquiries and frequently asked questions, Llama 3.1 allows human agents to focus on more complex issues that require a personal touch, thereby improving overall efficiency.

6.4. Personalized Interactions

- **Customization:** The model can be trained on specific business data, enabling it to provide personalized responses based on customer history and preferences. This personalization enhances the user experience by making interactions more relevant and tailored to individual needs.
- **Proactive Engagement:** Llama 3.1 can initiate proactive engagements, such as reminding customers of upcoming appointments, following up on recent purchases, or offering personalized recommendations based on past behavior.

6.5. Content Creation

Llama 3.1's prowess in generating high-quality, coherent, and contextually appropriate text makes it an invaluable tool for content creation. Key applications include:

6.5.1. Article Writing

- **High-Quality Output:** The model can draft articles on a wide range of topics, maintaining a high standard of quality and coherence. This capability is particularly useful for content marketing, journalism, and blogging, where timely and accurate content is essential.
- **Efficiency:** By automating the initial drafting process, Llama 3.1 allows writers to focus on refining and enhancing the content, significantly speeding up the content creation process.

6.5.2. Social Media

- **Engaging Posts:** The model can generate engaging social media posts tailored to different platforms, helping businesses and individuals maintain an active and appealing online presence. This includes crafting compelling captions, and hashtags and even responding to comments.
- **Consistency:** With Llama 3.1, content creators can ensure a consistent voice and style across all social media channels, reinforcing brand identity and messaging.

6.5.3. Creative Writing

- **Story Development:** The model can assist authors in developing story ideas, creating dialogues, and generating plot twists, making it a valuable tool for novelists, screenwriters, and other creative writers.
- **Idea Generation:** Llama 3.1 can provide inspiration and suggestions, helping writers overcome writer's block and explore new creative directions.

6.6. Language Learning

Llama 3.1's capabilities in language understanding and generation can significantly enhance language learning applications. Key benefits include:

6.6.1. Interactive Lessons

- **Engaging Content:** The model can generate interactive lessons and exercises that adapt to the learner's level and progress, making language learning more engaging and effective. This includes creating dialogues, quizzes, and comprehension exercises that challenge and motivate learners.
- **Real-Time Feedback:** Llama 3.1 can provide immediate feedback on exercises, helping learners understand their mistakes and improve their skills in real time.

6.6.2. Multilingual Support

- **Diverse Language Options:** The model supports multiple languages, making it suitable for learners of various linguistic backgrounds. This capability is particularly beneficial in multilingual classrooms or for self-learners exploring new languages.
- **Translation Practice:** Llama 3.1 can facilitate translation exercises, helping learners practice and refine their translation skills between different languages.

6.7. Personalized Learning

- **Tailored Exercises:** The model can create personalized learning paths and exercises based on the learner's strengths, weaknesses, and goals. This personalized approach ensures that learners receive the right level of challenge and support.
- **Cultural Context:** By incorporating cultural context into language lessons, Llama 3.1 helps learners understand the nuances and cultural aspects of the language they are studying, enhancing their overall proficiency and cultural awareness.

Llama 3.1's advanced NLP capabilities open up transformative opportunities in customer service, content creation, and language learning. Its ability to provide consistent, personalized, and high-quality text generation makes it an indispensable tool for businesses, educators, and content creators, driving efficiency, engagement, and innovation across these domains.

7. Limitations

Llama 3.1, as an advanced language model, represents a significant leap in natural language processing and machine learning. Its capabilities are impressive, but like all sophisticated

models, it is not without its limitations. Two notable limitations of Llama 3.1 are inherent bias and issues with interpretability.

7.1. Bias

Llama 3.1, like other large language models (LLMs), is susceptible to bias due to its extensive dataset sourced from various internet sources. This data can inadvertently learn and reproduce biases present in the training data, potentially reinforcing stereotypes or prejudices in sensitive areas like gender, race, and ethnicity. Additionally, Llama 3.1 might exhibit biases based on the distribution of topics in its training set, potentially leading to skewed or unbalanced coverage of certain subjects. Addressing these biases requires ongoing efforts to curate and refine training datasets and implement techniques to mitigate the influence of bias during the model's training and deployment phases. However, eliminating bias remains a challenging goal, and users must remain vigilant about the potential for biased outputs when interacting with the model.

7.2. Interpretability

Another significant limitation of Llama 3.1 is its interpretability. The model's decisions and outputs are often described as a "black box," meaning that understanding the exact reasons behind its responses can be challenging. Llama 3.1's architecture, which includes numerous layers and parameters, contributes to its powerful performance, but it also complicates the process of interpreting how the model arrives at a particular conclusion. Interpretability is crucial for several reasons. For instance, in applications requiring transparency—such as healthcare, legal, or financial sectors—understanding the rationale behind the model's outputs is essential for ensuring trust and accountability. If a model provides a recommendation or decision, stakeholders need to be able to trace the reasoning behind it to assess its validity and make informed decisions.

The complexity of Llama 3.1's neural network means that even experts may struggle to discern how individual inputs are transformed into specific outputs. This lack of transparency can hinder efforts to debug the model, improve its performance, and ensure that it operates within ethical and legal boundaries. Additionally, without clear interpretability, users may find it difficult to understand why the model behaves in certain ways, which can impact its adoption and integration into various applications. Efforts to enhance the interpretability of Llama 3.1 and similar models are ongoing. Researchers are exploring techniques such as model distillation, which involves creating simpler models that approximate the behavior of complex ones, and visualization tools that help users understand model behavior at a high level. However, achieving a balance between model complexity and interpretability remains an active area of research.

8. Conclusion

Llama 3.1 is a large language model that has achieved state-of-the-art results in natural language processing tasks. With a training dataset of 15 trillion tokens, it has been able to learn about language nuances and generate human-like text with remarkable accuracy. The diversity of training data, sourced from various domains, genres, and styles, has been critical to the model's success. Llama 3.1's capabilities have far-reaching implications for various industries and applications, including customer service, content creation, and language learning. It can power chatbots and virtual assistants, generate high-quality content, and aid language learning by generating interactive lessons, exercises, and conversations. Llama 3.1's ability to understand and generate human-like language has the potential to revolutionize the way we interact with technology. However, there is still much work to be done, and the model

will continue to refine and improve. In conclusion, Llama 3.1 is not just a model but a tool with the potential to transform the way we live, work, and communicate.

References

1. Meta AI. (2024). *Meta Llama 3.1*. Retrieved from <https://ai.meta.com/blog/meta-llama-3-1/>
2. Meta AI. (2024). *The LLaMA 3 Herd of Models*. Retrieved from <https://ai.meta.com/research/publications/the-llama-3-herd-of-models/>
3. Vavekanand, Raja and Kumar, Sanjay, LLMera: Impact of Large Language Models (June 05, 2024). Available at SSRN: <http://dx.doi.org/10.2139/ssrn.4857084>
4. Jiang, X., Yan, L., Vavekanand, R., & Hu, M. (2024). Large Language Models in Healthcare Current Development and Future Directions. <https://doi.org/10.20944/preprints202407.0923.v1>
5. Vavekanand, R., Karttunen, P., Xu, Y., Milani, S., & Li, H. (2024). Large Language Models in Healthcare Decision Support: A Review. <https://doi.org/10.20944/preprints202407.1842.v1>